

Data Science for Economists

Lecture 9: Intro to Numerical Methods

Alex Marsh

University of North Carolina | ECON 390

Table of contents

1. Introduction
2. Math and Stat Review
3. Intro to Numerical Methods

Introduction

Agenda

Today we will cover numerical methods and optimization.

This lecture might be a bit more theory heavy than programming heavy.

What is important here is to understand the concepts rather than the details.

- Unless you specialize in computational methods for economics later on, you will likely not need to understand the details.

Motivation

- There are many times there is an equation we would like to solve or know the value of.
- While ideally, we would solve these analytically to get an exact solution, for any non-trivial problem, this is either incredibly tedious or not possible.
- Today we will cover a handful of methods that you might need at some point.
- While some topics will relate within this lecture, most likely, they will feel disjoint and spattered.
- Understand that I am teaching a handful of tools that you will likely need at some point.

Introduction

Review

(Some important math and stat concepts)

Math Review: Derivative

- The derivative of a function tells you about it's rate of change.
 - "Instantaneous slope" (this is actually nonsensical the more one thinks about it)
- We denote the derivative of f with respect to x a few different ways:
 - $f'(x)$
 - $\frac{df}{dx}$
- The derivative tells us more than just the slope:
 - $f'(x) > 0 \rightarrow$ function is increasing
 - $f'(x) < 0 \rightarrow$ function is decreasing
 - $f'(x) = 0 \rightarrow$ function is not changing i.e. "critical point"
- The formal definition of a derivative is as follows:

$$f'(x) = \lim_{h \rightarrow 0} \frac{f(x+h) - f(x)}{h}$$

- We will actually be using this definition!

Math Review: Integrals

- Integrals calculate the area "under the curve"
- More broadly, they are used when you want to "sum up" a lot of very small things or take the average value of some function.
- The formal definition of the integral is

$$\int_a^b f(x)dx = \lim_{n \rightarrow \infty} \sum_{i=1}^n f(x_i) \Delta x$$

where $a \leq x_1 < x_2 < \dots < x_n \leq b$ and $x_{i+1} - x_i = \Delta x = (b - a)/n$

- We will be using this definition!

Intro to Numerical Methods

Approximating Derivatives

- While getting the analytical expression for $f'(x)$ is ideal due to accuracy and computational demand, sometimes it is not feasible to do so.
- Instead, we can approximate the derivative!
- Remember that

$$f'(x) = \lim_{h \rightarrow 0} \frac{f(x+h) - f(x)}{h}$$

- So for a "small" $h > 0$,

$$f'(x) \approx \frac{f(x+h) - f(x)}{h}$$

- This is the "forward differencing" approach; we can also "backward difference"

$$f'(x) \approx \frac{f(x) - f(x-h)}{h}$$

- These works, but we can get a more accurate approximation for a fixed h .

Approximating Derivatives (Cont.)

- Instead of forward or backward differencing, we can center difference:

$$f'(x) \approx \frac{f(x+h) - f(x-h)}{2h}$$

- Instead of favoring one side, we are approximating the derivative equally around a "neighborhood" of x .
- It is possible to show that the error generator by center differencing is smaller than right or left differencing.
- There are even more accurate methods that a bit more computationally demanding.
 - For those curious, look up Richardson's Extrapolation Method.

Example

Let's keep things simple and look at the derivative of $f(x) = x^2$. The analytical expression of the derivative is $f'(x) = 2x$.

```
f          = function(x){x^2}          #make function f(x)
fp         = function(x){2*x}         #make f'(x)
xs         = seq(0,2,0.5)              #store x's
h          = 0.01                      #store step size
der        = fp(xs)                    #calc actual derivatives
center_der = (f(xs+h)-f(xs-h))/(2*h)   #calc center differenced
forward_der = (f(xs+h)-f(xs))/(h)      #calc forward differenced
backward_der = (f(xs)-f(xs-h))/(h)     #calc backwards differenced
deriv_data = data.table("f'(x)"=der, "Center"=center_der,
                        "Foward"=forward_der, "Backward"=backward_der)

deriv_data
```

##	f'(x)	Center	Foward	Backward
## 1:	0	0	0.01	-0.01
## 2:	1	1	1.01	0.99
## 3:	2	2	2.01	1.99
## 4:	3	3	3.01	2.99
## 5:	4	4	4.01	3.99

Derivatives in Multiple Dimensions

- If instead of $f(x)$, we have something like $f(x_1, x_2, \dots, x_n)$, not much changes.
- The derivative will be a vector called the gradient denoted

$$\nabla_f = \left(\frac{\partial f}{\partial x_1}, \dots, \frac{\partial f}{\partial x_n} \right)$$

- To approximate each (partial) derivative, follow the same approach.
- The partial derivative of f in the x_i dimension at $x = (x_1, \dots, x_n)$ is,

$$\frac{\partial f}{\partial x_i} \approx \frac{f(x_1, \dots, x_i + h, \dots, x_n) - f(x_1, \dots, x_i - h, \dots, x_n)}{2h}$$

- Do this for all inputs and then collect them in a vector.
- This requires evaluating the function $2n$ times.

Example

Suppose that $f(x, y) = x^2 + \sin(y)$. $\frac{\partial f}{\partial x} = 2x$ and $\frac{\partial f}{\partial y} = \cos(y)$

```
f          = function(x){x[1]^2+sin(x[2])}      #make function f(x)
fp         = function(x){c(2*x[1],cos(x[2]))}    #make gradient of f(x)
h          = 0.01                             #store step size
fp(c(1,0.5))
```

```
## [1] 2.0000000 0.8775826
```

```
c((f(c(1+h,0.5))-f(c(1-h,0.5)))/(2*h),(f(c(1,0.5+h))-f(c(1,0.5-h)))/(2*h))
```

```
## [1] 2.0000000 0.8775679
```

```
c((f(c(1+h,0.5))-f(c(1,0.5)))/h,(f(c(1,0.5+h))-f(c(1,0.5)))/h)
```

```
## [1] 2.0100000 0.8751708
```

```
c((f(c(1,0.5))-f(c(1-h,0.5)))/h,(f(c(1,0.5))-f(c(1,0.5-h)))/h)
```

```
## [1] 1.9900000 0.879965
```

Approximating Integrals

- Just like derivatives, sometimes we might need to approximate an integral.
- Unlike derivatives, sometimes integrals don't have a closed-form expression at all!
- Therefore, all we can get for integrals is a numerical approximation.
- Remember that


$$\int_a^b f(x)dx = \lim_{n \rightarrow \infty} \sum_{i=1}^n f(x_i) \Delta x$$

- So for certain "partition" of points, $a \leq x_1 < x_2 < \dots < x_n \leq b$ and $x_{i+1} - x_i = \Delta x = (b - a)/n$,


$$\int_a^b f(x)dx \approx \sum_{i=1}^n f(x_i) \Delta x$$

- Note that a larger n makes this more accurate but means f must be evaluated n -times.
- Also note that depending on if we include a or b as a point makes this an upper sum or a lower sum.
 - Upper sums overestimate the integral.
 - Lower sums underestimate the integral.
 - Which to choose?


Upper sums and lower sums



Area of
region




Lower Sum



Upper Sum

The Trapezoidal Rule

- We don't need to pick upper or lower sums, we can use both! Sorta.
- Instead of rectangles, we use trapezoids.



The Trapezoidal Rule

- Each of the trapezoids is $\frac{a+b}{2}h$
- Here $h = \Delta x_i$, $a = f(x_{i-1})$, $b = f(x_i)$, $\Delta x_i = x_i - x_{i-1}$.
- So the area of each trapezoid is $\frac{f(x_{i-1})+f(x_i)}{2} \Delta x_i$
- Therefore, we know

$$\int_a^b f(x)dx \approx \sum_{i=1}^n \frac{f(x_{i-1}) + f(x_i)}{2} \Delta x_i$$

- The above expression works for an arbitrary grid i.e. Δx_i need not equal Δx_{i+k}
- If the grid is uniform, the above becomes a simple expression:

$$\sum_{i=1}^n \frac{f(x_{i-1}) + f(x_i)}{2} \Delta x_i = \frac{\Delta x}{2} (f(x_0) + 2f(x_1) + \dots + 2f(x_{n-1}) + f(x_n))$$

- This approximation is more accurate than the upper or lower sums.

Example

Suppose $f(x) = x^2$ and we want to evaluate $\int_0^{1.5} f(x)dx$.

```
f          = function(x){x^2}          #make function f(x)
Fx         = function(x){x^3/3}
trap_rule = function(func,a,b,n=20){
  dx = (b-a)/n
  xs = seq(a,b,dx)
  val = sum(2*func(xs[-c(1,n)]))
  val = val + func(xs[1])+func(xs[n])
  val*dx/2
}
true_val = Fx(1.5)-Fx(0)
my_approx_val20 = trap_rule(f,0,1.5,n=20)
my_approx_val100 = trap_rule(f,0,1.5,n=100)
approx_val = integrate(f,0,1.5)[[1]]
c(true_val,my_approx_val20,my_approx_val100,approx_val)
```

```
## [1] 1.125000 1.134633 1.125392 1.125000
```

Multivariate Integration

- Unlike derivatives, multidimensional integration becomes difficult quickly.
- One approach to multidimensional integration is actually Monte Carlo simulation, which we covered last lecture.
- A large class of multidimensional integration needed for probability distributions can be evaluated using Markov Chain Monte Carlo (MCMC) techniques.
 - Some examples include the Metropolis-Hastings algorithm and Gibbs Sampling.
 - These are needed for Bayesian inference.
 - Beyond the scope of this class, but worth looking into on your own.

Solving for Roots

- Many times you'll want to solve an equation numerically.
 - This may be because the closed-form solution is impossible or just tedious.
- When this is, we say we want to find the roots (i.e. x values) s.t.

$$f(x) = 0$$

- While we will cover the simplest methods to do so, the important thing to remember is the motivation and when you might want to use such techniques.
 - Better programmers than you have programmed faster and more robust methods.
- Dates back to the Babylonians and calculating square roots:

$$\sqrt{a} = x$$

or

$$x^2 - a = 0$$


Bisection Method

- One of the two methods we will cover here is the bisection method.
- While this method will always work, it can be slow.
- In order to use the method, you must have starting values a and b where

$$f(a) < 0 < f(b) \text{ or } f(a) > 0 > f(b)$$

- The idea is we take the midpoint between a and b , which we call c .
- If $\text{sign}(f(c)) = \text{sign}(f(a))$, replace a with c , otherwise b with c .
- Continue until $|f(c)| < \varepsilon$ where ε is a tolerance value.

Bisection Method



Reflection

- The bisection method works, but it is slow and sorta fancy "guess and check."
- We are only using the sign of the function to update the guesses.
- If we could incorporate more information of the function into the search process, it might be faster.
- Idea: Can we use the derivative of the function to tell us where to go next?
- This is known as Newton's method and it was the first improvement

Newton's Method

- Equation of tangent line in point-slope form at x :

$$f(x) - f(x_n) = f'(x_n)(x - x_n)$$

- Set $f(x) = 0$, replace x with x_{n+1} , and solve for x :

$$x_{n+1} = x_n - \frac{f(x_n)}{f'(x_n)}$$

- So make an initial guess x_0 , form $x_1 = x_0 - \frac{f(x_0)}{f'(x_0)}$.
- Keep updating x_{n+1} until $|f(x_{n+1})| < \varepsilon$
- Question: What technique that we learned earlier will most likely be needed for this method?

Comparing Methods

```
f = function(x){x^2-2}  
bisection_results = bisection(f,0,2)  
newton_results = Newtons(f,2)  
bisection_results
```

```
## [1] 1.414214 47.000000
```

```
newton_results
```

```
## [1] 1.414214 5.000000
```

Fixed Points

- Many times in economics, we are interested in finding a fixed point, which is defined as

$$f(x) = x$$

or

$$f(x) - x = 0$$

- Example: Supply and Demand equilibrium can be written as

$$P^S(Q^D(p)) = p$$

- At it's core, there is nothing different for solving fixed points.
 - Define $g(x) = f(x) - x$ and solve $g(x) = 0$.
- However, for some special functions, you can solve $f(x) = x$ via fixed point iteration.
- If the function f is a contraction, then we can simply iterate
 - $x_1 = f(x_0)$
 - $x_2 = f(x_1)$
 - $x_{n+1} = f(x_n)$
- What is a contraction?

Contractions

- Example: $f(x) = \cos(x)$

```
x0 = seq(0,2*pi,0.25)
xn = x0
for(n in 1:100){
  xnp1 = cos(xn)
  xn   = xnp1
}
xnp1
```

```
## [1] 0.7390851 0.7390851 0.7390851 0.7390851 0.7390851 0.7390851 0.7390851
## [8] 0.7390851 0.7390851 0.7390851 0.7390851 0.7390851 0.7390851 0.7390851
## [15] 0.7390851 0.7390851 0.7390851 0.7390851 0.7390851 0.7390851 0.7390851
## [22] 0.7390851 0.7390851 0.7390851 0.7390851 0.7390851
```

- Unfortunately, most functions are not contractions.
- However, there are some mathematical objects that are universally contractions and if you go onto graduate studies, you will experience them.
 - Dynamic programming and value functions.

Multivariate Root Finding

- You might need to find

$$f_i(x_1, \dots, x_n) = 0 \text{ for } i = 1, \dots, n$$

- There are methods similar to the ones we have already discussed for solving these problems.
- However, they are much outside the scope of this class.
- For now, it is sufficient to know that there exists a package in `R` for solving these problems: `nleqslv`.
- While the algorithms are more complex, the ideas remain the same.
- You will need `nleqslv` for the problem set.

Optimization

- The last numerical methods topic that we will cover is optimization.
- An optimization problem is one that can be written as follows:

$$\min_x f(x)$$

- Note, if we want to maximize, that's the same as minimizing $-f(x)$.
- We want to find the x value that minimizes $f(x)$.
- Sometimes, these will result in closed form solutions (e.g. OLS), but usually will be too complicated.
- We write the solution to this optimization problem as

$$x^* = \arg \min_x f(x)$$

- For functions that are differentiable, they are optimized at $f'(x) = 0$.
 - Look familiar?
- This is just root finding of $g(x) = 0$ where $f'(x) = g(x)$.

General Issues That Arise

- The first big issue that arises is that maxs, mins, and saddle points all have $f'(x) = 0$.
- Also, there can be local minimums if the function is not globally convex.
- This makes numerical optimization difficult!
- Using the second derivative/Hessian matrix can help solve some of these issues, but it is very computationally expensive to compute.
- As such, no one technique will always work.
- Lots of time and research goes into figuring out the best way to optimize one specific optimization problem.
 - E.g. check out this [best practices paper](#) for one model used in my subfield.
- Sometimes an entire project can live or die depending upon if you can optimize the objective function correctly.
 - E.g. one of my current projects...

Grid Search

- If the dimension of \mathbf{x} is small enough, you can make a grid of points to search on and see which set of values minimizes f .
- While this seems simple, in practice, you rarely want to do it.
- If \mathbf{x} has more than two dimensions, must search many points.
 - Particularly bad if f takes awhile to run.
- Must make grid fine enough so that you're not skipping over too many points, but too fine runs into the same problem as before.

```
x1 = seq(0,5,0.05)
x2 = seq(0,5,0.05)
x3 = seq(0,5,0.05)
nrow(expand.grid(x1,x2))
```

```
## [1] 10201
```

```
nrow(expand.grid(x1,x2,x3))
```

```
## [1] 1030301
```

Newton's Method... Again!

- We can actually use Newton's Method to solve optimization problems.
- Since we want to solve $f'(x) = 0$ and using $g(x) = f'(x)$, apply the formula from before:

$$x_{n+1} = x_n - \frac{g(x_n)}{g'(x_n)}$$

or

$$x_{n+1} = x_n - \frac{f'(x_n)}{f''(x_n)}$$

- Note that we must compute a derivative and a second derivative for each iteration.
- Computing the derivative requires evaluating the function twice.
- Computing the second derivative requires evaluating the function four times.
- Generally, if you have n inputs, you have to evaluate the function $4n^2/2$ times for the Hessian.
- So while this works, if your objective function takes awhile to run or has a lot of inputs, it might not be ideal.
 - Usually not ideal...

Gradient Descent

- All the second derivative does is scale the updating process so we don't "learn" too fast or too slow.
- Idea: Instead of calculating the scaling amount via the second derivative, we just set some parameter $\alpha > 0$ to update

$$x_{n+1} = x_n - \alpha f'(x_n)$$

- This is the idea behind gradient descent.
- Choosing α :
 - Sometimes α is just fixed at some small value.
 - However, you can also pick an optimal α :
 1. Calculate $f'(x_n)$ and save it.
 2. Then choose α to minimize $f(x_n - \alpha f'(x_n))$.
 3. Repeat this step each time you update x_{n+1} .
 - Whether this is beneficial is problem specific.
 - You have to solve a smaller optimization problem during your larger optimization problem.

Other Methods

- There are other methods.
 - Quasi-Newton: Calculate a function $B(x_n)$ that approximates the Hessian but is easier to calculate.
 - Derivative free methods:
 1. Nelder-Mead
 2. Simulated annealing
 3. BOBYQA, COBYLA
- Which method to choose?
 - Problem specific!
 - Depends on the properties of f , how long it takes f to run, how many inputs f has, etc.
 - Unfortunately, there is no one answer.
 - Optimization can make or break a project.

Other Methods

- Robustness: Performs well for various problems and starting values.
- Efficiency: Achieves the solution relatively quickly.
- Accuracy: Identify a solution with precision, not sensitive to starting values.

Up Next: Data Wrangling
